

Name: Nilesh Kawar


Roll No. 59

## Practical 2

Q1. Write a program to implement Linear Regression in python.

Program with Output:

```
#Importing Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
#Importing Dataset
data=pd.read_csv('linear regression.csv')
data.head()
```



The screenshot shows a Jupyter Notebook interface. The top part contains a code cell with the following Python code: `#Importing Libraries`, `import numpy as np`, `import pandas as pd`, `import matplotlib.pyplot as plt`, `#Importing Dataset`, `data=pd.read_csv('linear regression.csv')`, and `data.head()`. Below the code cell, the output is a preview of the first five rows of the 'linear regression.csv' dataset. The table has two columns: 'YearsExperience' and 'Salary'. The rows are indexed from 0 to 4.

	YearsExperience	Salary
0	1.1	39343
1	1.3	46205
2	1.5	37731
3	2.0	43525
4	2.2	39891

```
#Data processing
x=data.iloc[:, :-1].values      #independent variable array
y=data.iloc[:, 1].values        #dependent variable variable
                                vector
```

x



```
#Data processing
x=data.iloc[:, :-1].values      #independent variable array
y=data.iloc[:, 1].values        #dependent variable variable vector
```

x

```
array([[ 1.1],
       [ 1.3],
       [ 1.5],
       [ 2. ],
       [ 2.2],
       [ 2.9],
       [ 3. ],
       [ 3.2],
       [ 3.2],
       [ 3.7],
       [ 3.9],
       [ 4. ],
       [ 4. ],
       [ 4.1],
       [ 4.5],
       [ 4.9],
       [ 5.1],
       [ 5.3],
       [ 5.9],
       [ 6. ],
       [ 6.8]])
```



y

```
array([ 39343,  46205,  37731,  43525,  39891,  56642,  60150,  54445,
        64445,  57189,  63218,  55794,  56957,  57081,  61111,  67938,
        66029,  83088,  81363,  93940,  91738,  98273, 101302, 113812,
        109431, 105582, 116969, 112635, 122391, 121872])
```

```
#dataset into train and test data
```

```
from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split (x, y,
test_size=1/ 3, random_state = 0)
```

```
#Fitting Linear regression model into he training set
```

```
from sklearn.linear_model import LinearRegression
```

```
regressor=LinearRegression()
```

```
regressor.fit(x_train,y_train)
```

```

#dataset into train and test data
from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=1/3, random_state = 0)

[ ] #Fitting Linear regression model into the training set
from sklearn.linear_model import LinearRegression
regressor=LinearRegression()
regressor.fit(x_train,y_train)

LinearRegression()

```

```

#Predicting the test set results
y_pred = regressor.predict(x_test)
y_pred #predicted salaries

```

```

#Predicting the test set results
y_pred = regressor.predict(x_test)
y_pred #predicted salaries

array([ 40835.10590871, 123079.39940819,  65134.55626083,  63265.36777221,
        115602.64545369, 108125.8914992 , 116537.23969801,  64199.96201652,
        76349.68719258, 100649.1375447 ])

```

```

y_test #real salaries

```

```

y_test #real salaries

array([ 37731, 122391,  57081,  63218, 116969, 109431, 112635,  55794,
        83088, 101302])

```

```

#Visualizing the results

```

```

#plot for the train

```

```

plt.scatter(x_train, y_train, color="red")
    #plotting the observation line
plt.plot(x_train, regressor.predict(x_train), color='blue')
    #plotting the regression line
plt.title("Salary vs Experience (Training set)")
    #stating the title of the graph
plt.xlabel("Years of experience")
    #adding the name of x-axis
plt.ylabel("Salaries")
    #adding the name of y-axis
plt.show()
    #specifies end of graph

```



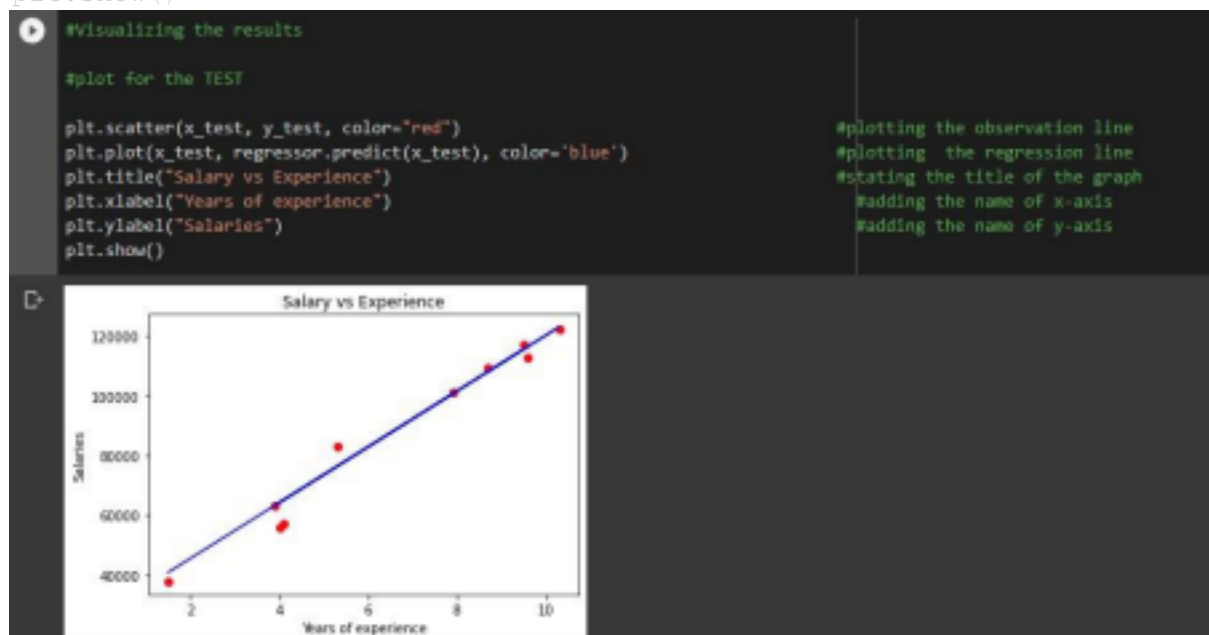
#Visualizing the results

#plot for the TEST

```

plt.scatter(x_test, y_test, color="red")
    #plotting the observation line
plt.plot(x_test, regressor.predict(x_test), color='blue')
    #plotting the regression line
plt.title("Salary vs Experience")
    #stating the title of the graph
plt.xlabel("Years of experience")
    #adding the name of x-axis
plt.ylabel("Salaries")
    #adding the name of y-axis
plt.show()

```



```
df=pd.DataFrame({'Actual' : y_test, 'Predicted' :
y_pred}) df
```

```
df=pd.DataFrame({'Actual' : y_test, 'Predicted' : y_pred})
df
```

	Actual	Predicted
0	37731	40835.105909
1	122391	123079.399408
2	57081	65134.556261
3	63218	63265.367772
4	116969	115602.645454
5	109431	108125.891499
6	112635	116537.239698
7	55794	64199.962017
8	83088	76349.687193
9	101302	100649.137545

**Q2. Write a program to implement Logistic Regression in python**

**Program with Output:**

```
#Importing Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
```

```
#importing dataset
data=pd.read_csv('logisticregression.csv')
data.shape
```

```
#Importing Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

#importing dataset
data=pd.read_csv('logisticregression.csv')
data.shape
```

(400, 5)

```
#display first 5 records
```

```
data.head()
```

```
#display first 5 records
data.head()
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
data.tail()
```

```
data.tail()
```

	User ID	Gender	Age	EstimatedSalary	Purchased
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

```
#info()->prints concise summary of dataframe
data.info()
```

```
#info()->prints concise summary of dataframe
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   User ID         400 non-null   int64
 1   Gender          400 non-null   object
 2   Age             400 non-null   int64
 3   EstimatedSalary 400 non-null   int64
 4   Purchased       400 non-null   int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
```

```
data.describe()
```

data.describe()

	User ID	Age	EstimatedSalary	Purchased
count	4.000000e+02	400.000000	400.000000	400.000000
mean	1.569154e+07	37.655000	69742.500000	0.357500
std	7.165832e+04	10.482877	34096.960282	0.479864
min	1.556669e+07	18.000000	15000.000000	0.000000
25%	1.562676e+07	29.750000	43000.000000	0.000000
50%	1.569434e+07	37.000000	70000.000000	0.000000
75%	1.575036e+07	46.000000	88000.000000	1.000000
max	1.581524e+07	60.000000	150000.000000	1.000000

```
#extracting dependent and independent variable
x=data.iloc[:,[2,3]].values #independent
y=data.iloc[:,-1].values #dependent
print(x)
```

```
#extracting dependent and independent variable
x=data.iloc[:,[2,3]].values #independent
y=data.iloc[:,-1].values #dependent
print(x)
```

```
[[ 19 19000]
 [ 35 20000]
 [ 26 43000]
 [ 27 57000]
 [ 19 76000]
 [ 27 58000]
 [ 27 84000]
 [ 32 150000]
 [ 25 33000]
 [ 35 65000]
 [ 26 80000]
 [ 26 52000]
 [ 20 86000]
 [ 32 18000]
 [ 18 82000]
 [ 29 80000]
 [ 47 25000]
 [ 45 26000]
 [ 46 28000]
 [ 48 29000]
 [ 45 22000]
 [ 47 49000]]
```

```
print(y)
```



```
#splitting the dataset into training and test set
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test =
train_test_split(x,y,test_size=0.30, random_state=0)
```

`x_test`



`x_train.shape,x_test.shape,y_train.shape,y_test.shape`





```
#feature Scaling
from sklearn.preprocessing import StandardScaler
st_x=StandardScaler()
x_train=st_x.fit_transform(x_train)
x_test=st_x.fit_transform(x_test)
```

```
print(x_train)
```



```
print(x_test)
```



```
#fitting logistic regression model to thr training
set from sklearn.linear_model import
LogisticRegression
Classifier=LogisticRegression(random_state=0)
Classifier.fit(x_train,y_train)
```



```
#predicting the test set result
y_pred=Classifier.predict(x_test)
y_trainPred=Classifier.predict(x_train)
print(y_pred)
print(y_trainPred)
```



```
#check model score
Classifier.score(x_test,y_test)
```



```
#creating the confusion matrix
from sklearn.metrics import confusion_matrix
cm=confusion_matrix(y_test,y_pred)
cm
```



```
#confusion matrix
print(confusion_matrix(y_train,y_trainPred))
```



```
#creating classification report
print(classification_report(y_test, y_pred))
```



**Q3. Write a python code to implement K-Nearest Neighbor(KNN) Algorithm for Machine Learning**

**Program with Output:**

```
#import Libraries
import numpy as np #perform a wide variety of mathematical operations
o n
import pandas as pd
from sklearn.metrics import confusion_matrix, classification_report

#importing dataset
data=pd.read_csv('logisticregression.csv')
data.head()
```



```
#extracting dependent and independent variable
x=data.iloc[:,[2,3]].values #independent variable array
y=data.iloc[:,4].values #dependent variable vector print(x)
```



```
print(y)
```



```
#splitting the dataset into training and testing set
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test =
train_test_split(x,y,test_size=0.30, random_state=0)
```

```
#feature Scaling
```

```
from sklearn.preprocessing import StandardScaler
st_x=StandardScaler()
x_train=st_x.fit_transform(x_train)
x_test=st_x.transform(x_test)
```

```
#fitting K-NN classifier to the training set
```

```
from sklearn.neighbors import KNeighborsClassifier
Classifier=KNeighborsClassifier(n_neighbors=5, metric='minkowski',
p=2) Classifier.fit(x_train,y_train)
```



```
y_pred=Classifier.predict(x_test)
print(y_pred)
```



```
#creating the confusion matrix
from sklearn.metrics import confusion_matrix
cm=confusion_matrix(y_test,y_pred)
cm
```



```
print(classification_report(y_test,y_pred))
```

